

An Explainable Artificial Intelligence Framework for Fraud Detection in Digital Banking Environments

Unza Munaf

*Department of Finance and Technology
FAST NUCES, Islamabad, Pakistan
i225502@isb.nu.edu.pk*

Abstract—Given the complexity of transactions and the sophistication of fraudsters, financial fraud poses serious threats to digital banking. The use of machine learning may detect fraud effectively but the existing systems function as black boxes. They provide the prediction but for that prediction, they do not explain it either. Because of that regulators cannot trust them. Moreover, organisations may also find them hard to trust. This work presents an XAI framework for fraud detection in digital banking environments. The system utilizes ensemble classifiers through gradient boosting along with SHAP and LIME which is able to effectively detect fraud and explain the reasoning. The system has a five-stage pipeline. Transactional data is pre-processed and features engineering is carried out and it generates enriched representation. In this, we extract temporal features, calculate the distance feature and card identification feature. A combined approach of XGBoost and LightGBM ensemble model for fraud classification with a dedicated explainability module generating per-prediction attribution scores using SHAP Shapley values and LIME local surrogates. The experiments yield an F1-Score of 0.9438 and an AUC-ROC of 0.9942 which is much higher than the set thresholds of 0.90 and 0.97. The model output is sufficiently explainable for fraud analysts, compliance teams and regulators to take action. In future work, we will adapt the framework to work with streaming and federated learning.

Index Terms—explainable artificial intelligence, fraud detection, digital banking, SHAP, XGBoost, ensemble learning

I. INTRODUCTION

The arrival of digital banking has revolutionized the financial services sector, allowing millions of users worldwide to transfer money, carry out transactions and manage accounts using mobile and web-based platforms. [9]. This shift toward digital-first banking has accelerated financial inclusion, reduced operational costs, and enhanced user convenience, making real-time payments and 24/7 account access the new standard rather than the exception.

Nevertheless, digitization of the banking sector has led to tremendous opportunities for financial frauds. The traditional fraud observed in brick-and-mortar establishments involved the physical forgery of documents, the impersonation of other parties, and other documents. In contrast, digital fraud occurs in relation to networked systems. This takes place through more sophisticated techniques such as card-not-present transactions on stolen credit cards, fraud through account takeover and synthetic identity creation. Another one is complex money laundering schemes. The Association of

Certified Fraud Examiners estimates that organizations lose approximately 5% of annual revenue to fraud, with digital banking institutions facing disproportionately higher risks due to transaction volume, velocity, and attack surface complexity [1].

Today systems that detect fraud mostly use machine learning. This includes things like neural networks and random forests and gradient boosting classifiers. These systems are really good at predicting things because they learn from transaction data. They can see patterns that're hard to understand. Fraud detection systems, like these are able to find fraud by looking at what happened in the past. They use this information to make guesses about what might happen in the future. Fraud detection systems and machine learning work well together to catch people who are trying to cheat the system. [10], [11]. These models excel at identifying subtle anomalies and non-linear relationships that rule-based systems cannot capture, leading to their widespread adoption across major financial institutions for real-time fraud screening and risk assessment.

Despite their strong performance, contemporary machine learning fraud detectors suffer from critical limitations that hinder their practical deployment. First, these models function as black boxes, providing predictions without human-interpretable justification for why specific transactions were flagged. Second, the lack of transparency creates regulatory compliance risks, particularly under frameworks like the EU AI Act and GDPR that mandate algorithmic accountability and the right to explanation [2], [3]. Third, fraud analysts cannot efficiently investigate flagged cases without understanding the model's reasoning, leading to operational bottlenecks and increased false positive costs.

These limitations are important because fraud detection is not a technical issue. It also involves laws and rules. Fraud detection affects how a company works and how customers are treated. The decisions made about fraud detection must be clear and justifiable. This is why these limitations matter in fraud detection. They help make sure that fraud detection is fair and responsible. [12]. Regulators require auditable evidence that automated systems do not exhibit discriminatory bias. Fraud investigation teams need interpretable signals to prioritize cases and gather supporting evidence. Customers

whose legitimate transactions are declined deserve comprehensible explanations rather than opaque algorithmic verdicts. The absence of explainability therefore represents both a compliance liability and a barrier to human-AI collaboration in fraud operations.

This paper addresses these challenges by proposing a hybrid Explainable Artificial Intelligence (XAI) framework that integrates high-performance ensemble learning with post-hoc interpretation techniques [5], [6]. The framework combines gradient boosting classifiers for fraud prediction with SHAP (SHapley Additive exPlanations) and LIME (Local Interpretable Model-agnostic Explanations) to generate per-prediction attribution scores quantifying each feature’s contribution to the final decision. By preserving predictive accuracy while adding transparency, the system enables regulatory-compliant, auditable, and operationally actionable fraud detection suitable for real-world digital banking deployment.

The remainder of this paper is organized as follows. Section II reviews related literature. Section III presents the problem statement. Section IV identifies the research gap and motivation. Section V describes the proposed methodology. Section VI covers the experiment design. Section VII details the implementation. Section VIII presents results. Section IX provides results analysis. Section X discusses limitations. Section XI concludes the paper.

A. Main Contributions and Novelty

The key contributions of this paper are as follows:

- **Propose** a unified XAI framework that jointly optimizes fraud detection accuracy and regulatory-compliant explainability, addressing a gap not resolved by any single prior approach [12].
- **Design** a dual-explainability module integrating SHAP Shapley value attribution and LIME local surrogate models to generate both global and instance-level explanations for every flagged transaction [5], [6].
- **Develop** a soft-voting ensemble combining XGBoost and LightGBM with SMOTE-balanced training, achieving F1-Score of 0.9438 and AUC-ROC of 0.9942, exceeding both target thresholds [7], [8].
- **Analyze** the impact of six domain-specific engineered features including geographic distance, night transaction flag, and age risk tier on fraud detection performance and SHAP attribution stability [5].
- **Demonstrate** regulatory alignment with GDPR Article 22 and the EU AI Act by producing auditable per-prediction explanation reports suitable for fraud analyst investigation and compliance audit trails [2], [3].

II. LITERATURE REVIEW

The integration of machine learning into fraud detection has significantly improved predictive capabilities; however, challenges related to model interpretability remain largely unresolved [12]. Existing research usually looks at accuracy and explainability as goals. It does not try to improve both at the time. In the world decisions affect customers and have to

follow rules. So it is very important to have performance and be transparent. This part of the work looks at what others have done in three areas: basic ways to detect fraud, explainable artificial intelligence and ways to combine learning methods. The financial domains need to have both performance and transparency because financial domains decisions directly impact customers and financial domains have to comply with regulatory standards. The section reviews work on foundational fraud detection approaches, explainable artificial intelligence and ensemble learning methods.

Early fraud detection studies primarily relied on statistical and rule-based techniques. Bolton and Hand [9] proposed a widely recognized taxonomy categorizing detection methods into supervised, unsupervised, and hybrid approaches. Supervised models learn from labeled transaction data, while unsupervised techniques identify anomalies based on deviations from normal behavior. Hybrid approaches combine both strategies to enhance detection performance. Although this framework laid the groundwork for subsequent research, it predates modern machine learning advancements and does not address issues such as scalability, model complexity, or interpretability, which are critical in contemporary systems.

The growing need for transparency in automated decision-making has led to the development of explainable AI methods. Ribeiro *et al.* [6] introduced LIME (Local Interpretable Model-agnostic Explanations), which provides local interpretability by approximating a model’s decision boundary using simple, interpretable models. By perturbing input features and observing prediction changes, LIME generates instance-specific explanations applicable to any classifier. Despite its flexibility, LIME suffers from instability, as explanations may vary for similar inputs, and it lacks global interpretability guarantees.

To overcome these limitations, Lundberg and Lee [5] proposed SHAP (SHapley Additive exPlanations), a framework based on cooperative game theory. SHAP assigns each feature a contribution value reflecting its impact on the model’s output, ensuring properties such as consistency and local accuracy. The TreeExplainer variant enables efficient computation for tree-based models, making SHAP particularly suitable for fraud detection tasks involving gradient boosting algorithms. Compared to LIME, SHAP provides more reliable and theoretically grounded explanations; however, it primarily focuses on feature-level importance and does not inherently capture temporal relationships in transaction data.

In parallel, ensemble learning techniques have become the dominant approach for fraud detection using structured datasets. Chen and Guestrin [7] introduced XGBoost, a scalable gradient boosting framework incorporating regularization, sparsity-aware split finding, and parallel processing. Its ability to handle missing values and prevent overfitting makes it highly effective for financial datasets, which are often noisy and complex. XGBoost has consistently demonstrated strong performance across various classification tasks, including fraud detection.

Building on this, Ke *et al.* [8] developed LightGBM, which enhances computational efficiency through histogram-based

learning, Gradient-based One-Side Sampling (GOSS), and Exclusive Feature Bundling (EFB). LightGBM employs a leaf-wise growth strategy that can achieve higher accuracy compared to traditional level-wise methods, particularly in high-dimensional data. However, this approach may increase the risk of overfitting if not properly controlled. Both XGBoost and LightGBM have achieved state-of-the-art results, motivating their use in ensemble-based fraud detection systems [7], [8].

Despite these advancements, notable research gaps remain. Arrieta *et al.* [12] highlighted that financial fraud detection requires both high predictive accuracy and interpretability, yet most existing approaches focus on only one of these aspects. High-performing models are often complex and lack transparency, while interpretable models tend to sacrifice predictive power. Additionally, many frameworks do not provide detailed, instance-level explanations required for regulatory compliance and auditing purposes [2], [3].

To deal with these problems people have started looking into ways that bring together ensemble learning and explainability techniques. This means that researchers are trying to find ways to make learning better by adding explainability techniques to it. The goal is to make ensemble learning more understandable and useful by using these explainability techniques. Ensemble learning is a way of combining models to get better results. By using explainability techniques with learning researchers hope to make it more effective and easy to understand. [5]–[7]. By leveraging the predictive strength of gradient boosting models alongside XAI methods such as SHAP and LIME, it is possible to achieve both accuracy and interpretability. Such integrated frameworks are essential for deploying reliable and transparent fraud detection systems in real-world financial environments.

Table I summarizes the reviewed literature and identifies the gaps addressed by the proposed framework.

III. PROBLEM STATEMENT

Given a stream of digital banking transactions [1]

$$\mathcal{T} = \{t_1, t_2, \dots, t_n\} \quad (1)$$

where each transaction t_i is characterized by a feature vector $\mathbf{x}_i \in \mathbb{R}^d$ and a binary fraud label $y_i \in \{0, 1\}$, the objective is to develop a real-time fraud detection system:

$$f : \mathbb{R}^d \rightarrow \{0, 1\} \times \mathcal{E} \quad (2)$$

that simultaneously produces accurate fraud predictions and human-interpretable explanations.

The system must maintain sub-100 ms prediction latency for real-time transaction authorization while operating on feature-sparse datasets typical of digital banking environments [7], [8]. Regulatory frameworks, including GDPR Article 22 and the EU AI Act [2], [3], mandate algorithmic transparency and auditable decision rationales for automated decisions affecting individuals. Financial institutions must therefore provide comprehensible explanations to customers whose transactions are flagged, both to satisfy regulatory obligations and to maintain trust in digital services.

From an operational perspective, fraud analysts require interpretable model outputs to efficiently prioritize investigations and distinguish genuine threats from false positives across large volumes of daily alerts [12]. When a lot of transactions are marked as fake people who study these transactions need to find out what is causing these suspicions. They have to do this and they cannot look at every single transaction that happened before. If we do not know why something is happening it takes a time to look into it. This causes delays. The company has to spend more money to fix the problem. The people studying the transactions need to know what is going on with the transactions that are flagged as fake.

Existing approaches present a fundamental trade-off between predictive performance and interpretability [9], [10]. Black-box models like neural networks and gradient boosting classifiers are really good at detecting things. They get it most of the time.. We do not know how they work. Black-box models like neural networks and gradient boosting classifiers are not transparent. On the hand models like logistic regression are easy to understand. We can see how logistic regression models work. Logistic regression models are not as good at detecting things as black-box models, like deep neural networks and gradient boosting classifiers are.

This research addresses the challenge of preserving state-of-the-art fraud detection performance while ensuring interpretability [5], [6]. The goal is to get really good at predicting things. We want to achieve scores like an F1-Score of more than 0.90 and an AUC-ROC of more than 0.97. At the time we need to make sure our explanations are trustworthy and easy to understand. This is important for banking because we have to follow rules and be transparent about how our systems work. We want to make sure our systems are good enough to use in life in production digital banking environments and that they meet all the necessary requirements for regulatory compliance and operational transparency which is crucial, for digital banking. The digital banking environments require that we have predictive metrics, including F1-Score and AUC-ROC and that we provide faithful explanations.

IV. RESEARCH GAP AND MOTIVATION

Despite significant advances in applying machine learning to fraud detection, existing research exhibits critical limitations that prevent effective real-world deployment [9], [12]. The most prominent gap is the near-universal trade-off between predictive performance and interpretability. High-performing models including deep neural networks, XGBoost, and LightGBM consistently outperform interpretable alternatives such as logistic regression and decision trees, but operate as opaque black boxes providing no insight into individual prediction reasoning [7], [8].

Regulatory frameworks including the EU General Data Protection Regulation (GDPR) Article 22 and the EU AI Act [2], [3] explicitly mandate that automated decision systems affecting individuals be accompanied by meaningful explanations. Existing fraud detection literature rarely addresses this compliance dimension, creating a substantial gap between

TABLE I: Recent Proposed Models in Literature with Possible Limitations

Ref No.	Proposed Model	Tools/Technology	Experimentation	Contributions	Limitations
[1]	AI-Powered Fraud Detection System	LightGBM, Rule-based systems	Comparative analysis on transaction datasets	Proves AI models detect 35% more fraud than rule-based systems	Limited to a single model type; no explainability layer
[2]	AI-Driven Fraud Detection via GNNs	Graph Neural Networks (GNN), Deep Learning	Evaluation on money laundering datasets	Identifies GNNs as superior for detecting complex laundering schemes	High computational cost; lacks interpretability for analysts
[3]	Ensemble ML for Transactions	Random Forest, XGBoost, Scikit-learn	Benchmarked on financial transaction datasets	Highlights ensemble methods for scalability in fraud detection	No explainability; limited to structured tabular data
[4]	Operational Capability Triad (ORT)	ML governance frameworks, audit tools	Conceptual and governance-level analysis	Introduces ORT for model governance and compliance tracking	Theoretical; lacks experimental validation on real fraud data
[5]	Decentralized ML Fraud Framework	Federated Learning, GDPR-compliant tools	Evaluated in simulated online banking environments	Proposes decentralized ML to comply with GDPR regulations	Performance degrades in non-IID data distributions
[6]	XAI-Augmented Fraud Classifier	SHAP, LIME, Black-box ML models	Applied to auditable fraud classification tasks	Incorporates SHAP/LIME layers to make black-box AI auditable	Explanation stability not rigorously evaluated
[7]	Stacking Ensemble Fraud Detector	XGBoost, CatBoost, Stacking meta-learner	Tested on credit card fraud benchmark datasets	Proposes stacking ensemble of XGBoost and CatBoost for precision	No instance-level explanation generation included
[8]	XAI-Driven Banking Transparency Model	XAI frameworks, Trust metrics	Analytical review of explainability in banking	Analyzes how XAI bridges algorithmic clarity and institutional trust	No quantitative fraud detection benchmarks provided
[9]	Feature Engineering for Imbalanced Fraud	SMOTE, Random Forest, XGBoost	Experimented on imbalanced banking datasets	Investigates feature engineering impact on imbalanced datasets	Does not address post-hoc explainability requirements
[10]	Blockchain-AI Transparent Framework	Blockchain, AI, IEEE Access protocols	Comparative review of privacy-preserving AI	Reviews integration of blockchain and AI for secure data sharing	Practical deployment complexity not addressed
[11]	Real-Time Credit Card Fraud Detector	Deep Learning, High-dimensional feature spaces	Evaluated on high-volume real-time transaction streams	Focuses on high-dimensional data and real-time processing limits	Scalability issues; no interpretability mechanisms
[12]	XAI for Mobile Wallet FinTech	SHAP, Mobile AI, Digital wallet APIs	Applied to real-time mobile transaction decisions	Examines real-time decision transparency for mobile digital wallets	Limited to mobile domain; not generalized to banking
[1], [2]	GNN with Lambda Architecture	GNN, Lambda Architecture, Graph databases	Tested on dynamic transaction flow datasets	Proposes GNN with lambda architecture for dynamic transaction flows	High latency; explanation layer absent

academic performance benchmarks and real-world regulatory requirements. Financial institutions deploying black-box models face legal liability exposure and regulatory penalties despite achieving high detection accuracy.

Furthermore, extreme class imbalance where fraudulent transactions constitute 0.1% to 3.5% of total volume remains inadequately handled [11]. Although techniques such as SMOTE and cost-sensitive learning have been explored individually, their systematic integration into end-to-end explainable pipelines is underrepresented. Most studies applying SMOTE fail to assess whether synthetic samples distort feature attributions produced by post-hoc explanation methods like SHAP or LIME [5], [6].

Additionally, existing XAI-based fraud detection studies

evaluate explanations only qualitatively or in controlled settings, without assessing explanation fidelity, stability, or operational usability from fraud analysts' perspectives [12]. We really need a system that combines machine learning, fair handling of uneven data and clear explanations for how it works so it can be used in real digital banking. This research is trying to fill that gap by creating a framework that does all these things. Digital banking needs a system like this to work properly. The system should have machine learning and fair handling of uneven data and it should be able to explain what it is doing so we can trust it. This is what this research is trying to achieve with its framework, for digital banking.

V. PROPOSED METHODOLOGY

The proposed framework implements a five-stage pipeline transforming raw transaction data into explainable fraud predictions: *Data Collection*, *Data Preprocessing*, *Feature Engineering*, *Ensemble Model Design*, and *XAI-Driven Prediction System* [12]. Each stage addresses specific challenges including feature sparsity, overfitting risk in limited feature spaces, and regulatory interpretability requirements. Figure 1 illustrates the complete system architecture and information flow across pipeline stages.

A. Data Collection

The framework utilizes the Kaggle Credit Card Fraud Detection dataset comprising 10,000 sampled real-world transaction records [1]. The dataset contains 23 raw attributes including transaction timestamp, credit card number, merchant details, transaction amount, cardholder demographics (gender, job, date of birth), geographic coordinates, and a binary fraud label (*is_fraud*). From the full dataset of 555,719 transactions, 2,145 confirmed fraud cases were combined with 7,855 randomly sampled legitimate transactions to form the experimental corpus, preserving authentic fraud signal distribution.

B. Data Preprocessing

The preprocessing pipeline prepares raw features for model training while maintaining data integrity [11]. Irrelevant columns (first name, last name, street address, transaction number) are dropped. Transaction timestamps are parsed to extract hour, day-of-week, and month components. Date-of-birth fields are converted to cardholder age in years relative to the dataset reference date.

Categorical features including merchant category, gender, city, state, and job title are converted to numerical representations using label encoding. Numerical features including transaction amount, geographic coordinates, city population, merchant coordinates, and derived age are normalized using robust scaling to minimize the influence of outliers. The dataset contains no missing values, eliminating the need for imputation.

C. Feature Engineering

Feature engineering constructs six domain-specific predictors from the base attributes to enrich model input with behavioral, spatial, and temporal fraud signals [9]:

- **Geographic Distance:** Euclidean distance between cardholder coordinates and merchant coordinates, capturing spatial anomalies indicative of card-not-present fraud.
- **Amount Quantile Bin:** Transaction amount discretized into four quartile bins, encoding relative spending magnitude.
- **Night Transaction Flag:** Binary indicator for transactions occurring between 00:00 and 05:00, a high-risk temporal window.
- **Weekend Flag:** Binary indicator for Saturday or Sunday transactions.

- **Card Last-4 Digits:** Final four digits of the credit card number extracted as an identifier integrity signal.
- **Age Risk Tier:** Cardholder age segmented into four risk tiers (under 25, 26–40, 41–60, over 60).

These engineered features expand the input representation from 17 preprocessed features to 20 total features entering the ensemble model.

D. Ensemble Model Architecture

The detection engine employs gradient boosting using XGBoost [7] and LightGBM [8] as base learners. XGBoost provides strong regularization, while LightGBM ensures computational efficiency.

Models are configured with:

- Maximum depth = 6
- Minimum child weight = 5
- Learning rate = 0.05
- L2 regularization = 1.0
- Number of estimators = 200

Training uses an 80–20 stratified split. SMOTE oversampling is applied exclusively to the training partition to balance fraud and legitimate classes before model fitting. Predictions are aggregated via soft voting:

$$p_{\text{ensemble}} = \frac{p_{\text{XGBoost}} + p_{\text{LightGBM}}}{2} \quad (3)$$

Hyperparameters are optimized using grid search with five-fold cross-validation, maximizing F1-score [11].

E. XAI Explainability Module

The explainability layer integrates SHAP [5] and LIME [6]. SHAP computes feature contributions using Shapley values. For a transaction $\mathbf{x}$:

$$f(\mathbf{x}) = \phi_0 + \sum_{i=1}^d \phi_i \quad (4)$$

where ϕ_0 is the base value and ϕ_i represents the contribution of feature i . Positive values increase fraud probability, while negative values decrease it.

LIME constructs local surrogate models to approximate the ensemble's decision boundary for individual flagged transactions [6].

The module produces:

- Global feature importance rankings
- Per-transaction SHAP waterfall explanations
- LIME local linear surrogate explanations

These outputs support regulatory compliance such as GDPR Article 22 [3].

F. Decision System and Operational Integration

Each transaction receives a fraud probability score $p \in [0, 1]$ [12]. A threshold $\tau = 0.5$ is applied for binary classification:

$$\hat{y} = \begin{cases} 1 \text{ (fraud)} & \text{if } p \geq \tau \\ 0 \text{ (legitimate)} & \text{if } p < \tau \end{cases} \quad (5)$$

Fraudulent transactions are flagged and routed to case management systems with SHAP and LIME explanation reports attached. Explanation outputs are stored in audit logs to satisfy compliance requirements such as PCI-DSS [4] and EU AI Act transparency obligations [2], [3].




Fig. 1: System Architecture of the Proposed XAI Fraud Detection Framework.

VI. EXPERIMENT DESIGN

The experimental design is structured to rigorously evaluate both the predictive performance and the explainability quality of the proposed XAI framework [12]. Experiments are conducted in three phases: baseline model benchmarking, ensemble model evaluation, and XAI explanation quality assessment. All experiments use the Kaggle Credit Card Fraud Detection dataset with a stratified 80/20 train-test split. Five-fold stratified cross-validation on the training set is used for hyperparameter optimisation. Baseline comparisons include Logistic Regression, Decision Tree, and Random Forest classifiers trained on the same features to contextualise the performance of the proposed ensemble [10], [11].




Fig. 2: System Workflow of the Proposed XAI Fraud Detection Framework.

A. Dataset Description

The experimental evaluation employs a sampled credit card fraud dataset comprising 10,000 transaction records drawn from the Kaggle fraudTest.csv corpus (555,719 total records) [1]. The sample includes all 2,145 confirmed fraud transactions supplemented with 7,855 randomly selected legitimate transactions. Table II presents the complete dataset characteristics.

TABLE II: Dataset Description

Attribute	Details
Dataset Source	Kaggle Credit Card Fraud Detection
Total Transactions	10,000 records (sampled)
Fraud Cases	2,145 (21.45%)
Legitimate Cases	7,855 (78.55%)
Number of Features	20 (after preprocessing & engineering)
Feature Types	Categorical: category, gender, job, city, state Numerical: amt, city_pop, age Temporal: trans_hour, trans_day, trans_month Engineered: geo_distance, night_txn, amt_bin
Missing Values	None (complete records)
Class Imbalance	Handled via SMOTE on training set
Data Split	80% Training (8,000) / 20% Testing (2,000) Stratified to preserve class proportions
After SMOTE	Training: 6,284 Legit 6,284 Fraud

B. Tools and Technologies

Table III summarizes the complete technology stack employed for framework implementation. All experiments were conducted in Google Colaboratory using Python 3.10.

SHAP and LIME were selected as the explainability backbone due to their model-agnostic design, extensive validation in financial fraud literature, and complementary global-local explanation capabilities [12]. SHAP provides theoretically grounded Shapley values satisfying consistency guarantees [5], while LIME offers intuitive local linear approximations accessible to non-technical fraud analysts [6].

TABLE III: Tools and Technologies

Category	Tool/Library	Purpose
Language	Python 3.10	Core development
ML Frameworks	Scikit-learn 1.3	Baseline models
	XGBoost 2.0	Gradient boosting
	LightGBM 4.0	Gradient boosting
XAI Libraries	SHAP 0.43	Shapley attribution
	LIME 0.2	Local explanations
Imbalance	imbalanced-learn	SMOTE oversampling
Data Processing	Pandas 2.1	Data manipulation
	NumPy 1.25	Numerical arrays
Visualization	Matplotlib 3.8	Static plots
	Seaborn 0.13	Statistical viz
Development	Google Colab	Cloud environment

C. Evaluation Metrics

Given the class imbalance in the experimental dataset, multiple complementary metrics are employed to provide a comprehensive view of model performance [9]. Table IV presents the complete evaluation framework.

TABLE IV: Evaluation Metrics and Justification

Metric	Formula	Relevance
Accuracy	$\frac{TP+TN}{TP+TN+FP+FN}$	Overall correctness
Precision	$\frac{TP}{TP+FP}$	Minimizes false alerts
Recall	$\frac{TP}{TP+FN}$	Catches actual fraud
F1-Score	$\frac{2 \cdot P \cdot R}{P+R}$	Balances P and R
AUC-ROC	Area under ROC	Discriminative power
MCC	Correlation coef.	Balanced class metric
AUPRC	Area under PR curve	Positive class focus

Target: $F1\text{-Score} > 0.90$, $AUC\text{-ROC} > 0.97$

Accuracy alone is insufficient for imbalanced fraud datasets [9]. Precision measures the fraction of flagged transactions that are genuinely fraudulent, directly controlling false alert volume. Recall captures the fraction of actual fraud cases detected, critical for minimizing financial loss. F1-Score harmonizes precision and recall into a single metric. AUC-ROC evaluates discriminative power across all decision thresholds. MCC provides a balanced measure robust to class imbalance. AUPRC is particularly informative for imbalanced positive-class detection tasks [11].

VII. IMPLEMENTATION

The complete framework was implemented in Python 3.10 within the Google Colaboratory cloud environment, enabling reproducible execution without local hardware dependencies [11]. The implementation follows the five-stage pipeline defined in the proposed methodology.

A. Development Environment

The system was developed and executed in Google Colaboratory using a CPU runtime [11]. All library dependencies were installed via `pip`, and the dataset was uploaded directly to the Colab session filesystem. The development environment is fully reproducible: all random seeds are fixed to 42 across NumPy, Scikit-learn [11], XGBoost [7], and LightGBM [8] to ensure deterministic results.

B. Data Pipeline Implementation

The data pipeline was implemented using Pandas for tabular manipulation and NumPy for numerical operations [11]. Raw CSV data was loaded and a stratified sample of 10,000 records was extracted, preserving all 2,145 fraud cases. Datetime parsing extracted three temporal features from transaction timestamps. Label encoding was applied to five categorical columns using Scikit-learn's `LabelEncoder`. Robust scaling was applied to seven numerical columns using `RobustScaler`, which centers and scales by the interquartile range, providing resilience against transaction amount outliers.

C. Model Training

All models were trained on SMOTE-balanced training data. The XGBoost [7] and LightGBM [8] models were configured with identical hyperparameters: maximum depth of 6, minimum child weight/samples of 5, learning rate of 0.05, L2 regularization of 1.0, and 200 estimators. Baseline models (Logistic Regression, Decision Tree, Random Forest [10]) were trained with default Scikit-learn [11] configurations except for a Decision Tree maximum depth of 6 and Random Forest with 100 estimators. All models received the same SMOTE-balanced training set to ensure a fair comparison.

D. Explainability Implementation

SHAP explanations were generated using `shap.TreeExplainer` applied to the trained XGBoost model [5]. Global feature importance was visualized using SHAP bar plots aggregating mean absolute Shapley values across the test set. Instance-level explanations were generated using SHAP beeswarm plots showing the directional impact of each feature. LIME explanations were generated using `LimeTabularExplainer` with 10 features selected per instance, producing local linear surrogate models for individual fraud-flagged transactions [6].

TABLE V: Model Performance Comparison on Test Set (2,000 Transactions)

Model	Accuracy	Precision	Recall	F1-Score	AUC-ROC	MCC	AUPRC
Logistic Regression	0.7855	0.0000	0.0000	0.0000	0.5279	0.0000	0.2321
Decision Tree	0.9345	0.8239	0.8834	0.8526	0.9547	0.8114	0.8743
Random Forest	0.9585	0.9240	0.8788	0.9008	0.9892	0.8751	0.9690
XGBoost	0.9750	0.9438	0.9394	0.9416	0.9934	0.9257	0.9813
LightGBM	0.9755	0.9502	0.9347	0.9424	0.9944	0.9269	0.9817
Ensemble (XGB+LGBM)	0.9760	0.9482	0.9394	0.9438	0.9942	0.9286	0.9826

Bold indicates best performance. Target thresholds: F1-Score > 0.90, AUC-ROC > 0.97 — both achieved by ensemble.

VIII. RESULTS

A. Model Performance Comparison

Table V presents the complete performance metrics for all six models evaluated on the held-out test set of 2,000 transactions [11].

The XGB+LightGBM ensemble achieves the highest overall performance with an F1-Score of 0.9438 and AUC-ROC of 0.9942, meeting both target thresholds defined in the problem statement [7], [8]. Logistic Regression fails entirely on the fraud class ($F1 = 0.0000$), confirming that linear models cannot capture the non-linear patterns present in transaction data without further feature transformation [9].

B. Confusion Matrix

Figure 3 presents the confusion matrix for the ensemble model on the test set [7], [8]. Out of 429 total fraud transactions in the test set, 403 were correctly identified (True Positives), yielding a recall of 93.94%. Only 26 fraud cases were missed (False Negatives). Of 1,571 legitimate transactions, 1,549 were correctly classified (True Negatives), with only 22 falsely flagged as fraud (False Positives), demonstrating strong precision of 94.82%.




Fig. 3: Confusion Matrix of the XGB+LightGBM Ensemble Model. TP=403, FP=22, FN=26, TN=1549.

C. Performance Metrics Comparison

Figure 4 illustrates the performance metrics across all six models [7], [8], [10]. The ensemble model consistently leads across all metrics. The clear performance gap between Logistic Regression and tree-based models confirms the non-linear nature of fraud patterns in this dataset.




Fig. 4: Performance Metrics Comparison Across All Models.

D. ROC Curves

Figure 5 presents ROC curves for all models [9], [11]. The ensemble achieves an AUC of 0.994, closely followed by LightGBM (0.994) and XGBoost (0.993). The Logistic Regression curve nearly coincides with the random classifier diagonal ($AUC = 0.528$), confirming its inability to discriminate fraud from legitimate transactions. All tree-based models exhibit steep initial rises, indicating high true positive rates at low false positive rate thresholds.

E. Precision-Recall Curves

Figure 6 presents Precision-Recall curves, particularly informative for imbalanced detection tasks [9], [11]. The ensemble achieves AUPRC of 0.983, maintaining near-perfect precision across a wide recall range. The Decision Tree curve degrades more rapidly at high recall, indicating increased false positives when the threshold is lowered. Logistic Regression achieves only $AUPRC = 0.232$, consistent with its failure on the fraud class.

F. SHAP Global Feature Importance

Figure 7 presents the SHAP global feature importance ranking derived from the XGBoost model [5]. Transaction amount (*amt*) is the dominant predictor with a mean absolute SHAP value of approximately 2.6, nearly double the second-ranked feature *amt_bin* (1.47). Merchant category (*category*) ranks third, followed by transaction hour (*trans_hour*).




Fig. 5: ROC Curves for All Models. The ensemble achieves AUC-ROC = 0.994.




Fig. 6: Precision-Recall Curves for All Models. Ensemble AUPRC = 0.983.

These findings confirm that fraud is primarily driven by anomalous transaction amounts and spending categories rather than geographic or demographic features.

G. SHAP Beeswarm Plot

Figure 8 presents the SHAP beeswarm plot showing the directional impact of each feature [5]. High values of *amt* (red dots) concentrate on the positive SHAP axis, confirming that high transaction amounts strongly increase fraud probability. Low values of *amt* (blue dots) have negative SHAP values, pushing predictions toward legitimate. The *amt_bin* feature shows a bimodal pattern. *trans_hour* shows that nighttime transactions (high values) increase fraud probability, consistent with the engineered *night_txn* feature.

H. LIME Local Explanation

Figure 9 presents the LIME local explanation for a single flagged fraud transaction [6]. The green bars indicate features




Fig. 7: SHAP Global Feature Importance (XGBoost). *amt* is the dominant fraud predictor.

supporting the fraud classification, while the red bars indicate features opposing it. For this transaction, *amt* in the range 0.42 to 4.05 provides the strongest positive contribution toward fraud classification. The *amt_bin* condition (2.00 to 3.00) acts as the strongest opposing signal. This instance-level explanation demonstrates that the framework satisfies GDPR Article 22 requirements [3] by providing comprehensible, actionable rationales for individual fraud decisions.

I. SMOTE Class Balancing

Figure 10 illustrates the effect of SMOTE oversampling on the training set [11]. Before SMOTE, the training partition contained 6,284 legitimate and 1,716 fraud records. After SMOTE, both classes are balanced at 6,284 samples each. This balancing is applied exclusively to the training set, ensuring the test set preserves the authentic class distribution for unbiased evaluation.

IX. RESULTS ANALYSIS

A. Proposed Method vs. Baseline Performance

The experimental results confirm that the proposed XGB+LightGBM ensemble substantially outperforms all baseline methods across every metric [7], [8]. Logistic Regression achieves 78.55% accuracy but completely fails to detect any fraud (F1 = 0.0000, Precision = 0.0000, Recall = 0.0000), demonstrating that linear decision boundaries are fundamentally inadequate for non-linear fraud patterns even




Fig. 8: SHAP Beeswarm Plot showing directional feature impact on fraud predictions.




Fig. 9: LIME Local Explanation for a flagged fraud transaction. Green bars support fraud classification; red bars oppose it.




Fig. 10: Effect of SMOTE on Training Set Class Distribution. Fraud samples increased from 1,716 to 6,284.

after SMOTE balancing. Its AUC-ROC of 0.5279 is barely above random classification [9], [11].

Decision Tree achieves competitive performance ($F1 = 0.8526$, $AUC-ROC = 0.9547$) but exhibits instability at high recall thresholds visible in its Precision-Recall curve [10]. Random Forest improves substantially ($F1 = 0.9008$, $AUC-ROC = 0.9892$), meeting the $F1$ threshold but falling short of the AUC-ROC target of 0.97 at the prediction level [10]. XGBoost [7] ($F1 = 0.9416$) and LightGBM [8] ($F1 = 0.9424$) individually surpass all baselines, while the soft-voting ensemble achieves the best overall $F1$ -Score of 0.9438 and AUC-ROC of 0.9942, exceeding both target thresholds.

B. Insights from SHAP Explainability

SHAP analysis [5] reveals that transaction amount is the single most important fraud signal, contributing approximately 35% of total mean absolute feature importance. This is consistent with real-world fraud patterns where fraudsters execute atypically large transactions to maximize returns before detection [1]. Merchant category ranks third, suggesting that fraud concentrates in specific spending categories. Transaction hour's presence in the top four features validates the engineered `night_txn` flag as capturing genuine temporal risk.

Geographic features (`lat`, `long`, `geo_distance`, `merch_lat`, `merch_long`) rank in the bottom quartile, suggesting that card-present fraud is less geographically anomalous than card-not-present schemes in this dataset [12]. Cardholder demographics (`age`, `gender`) contribute marginally, confirming that fraud targeting is not strongly demographic in this corpus.

C. Interpretation of LIME Explanations

The LIME local explanation [6] for the sample fraud transaction demonstrates the framework's operational utility. The explanation identifies specific feature conditions driving the fraud classification in human-readable form (e.g., “`amt` between 0.42 and 4.05 increases fraud probability”). Fraud analysts can directly act on such explanations by verifying the flagged transaction amount and category against the cardholder's historical profile, significantly reducing investigation time compared to opaque model outputs [12].

D. SMOTE Effectiveness

The SMOTE balancing step proved critical: without it, Logistic Regression completely failed on the minority class,

and tree-based models would likely exhibit recall bias toward the majority class [9], [11]. By generating synthetic minority samples in feature space, SMOTE enabled all tree-based models to learn robust fraud decision boundaries. The fact that the test set was kept unmodified ensures that reported metrics reflect realistic deployment performance.

X. LIMITATIONS

A. Dataset Constraints

The experimental evaluation uses a 10,000-record sample from a larger Kaggle dataset rather than the complete corpus of 555,719 transactions [1]. The sample has all the confirmed fraud cases which's 2,145.. The problem is the sample is not big enough. This means the model does not see many of the rare fraud patterns. To use this in the world we need to train the model on all the data or we need to keep updating it with new transactions. This way the model can see all the fraud patterns, including the ones.

The dataset originates from a single source reflecting a specific time period (2020) and geographic region [1]. Fraud patterns evolve rapidly as fraudsters adapt to detection systems, meaning model performance may degrade over time without periodic retraining on current transaction data [9].

B. Feature Availability

Several high-value fraud indicators commonly used in production systems are absent from this dataset, including device fingerprinting, IP geolocation, browser metadata, velocity features (number of transactions per hour), and merchant risk scores [12]. The absence of these features means the reported performance likely underestimates what a production-grade system would achieve with richer feature sets.

C. Model and Computational Constraints

SHAP computation scales quadratically with the number of features and samples; on large datasets this becomes computationally prohibitive without approximation methods such as KernelSHAP sampling [5]. LIME explanation stability is inherently limited by its perturbation-based approximation, meaning explanations for similar transactions may vary [6]. This instability could reduce fraud analyst trust in instance-level explanations.

The framework was evaluated in a batch processing context [7], [8]. Real-time deployment requiring sub-100ms per-transaction scoring with concurrent SHAP explanation generation would demand significant engineering optimization, including model serialization, explanation caching, and potentially approximated SHAP methods.

D. Regulatory and Ethical Scope

While the framework generates SHAP and LIME explanations to support GDPR Article 22 compliance [3], the regulatory adequacy of these explanations has not been formally assessed by legal or compliance experts [2]. The comprehensibility of SHAP Shapley values to non-technical customers (as required by GDPR's right-to-explanation provision) remains an open question requiring user studies [12].

XI. CONCLUSION

This paper is about a kind of Artificial Intelligence called Explainable Artificial Intelligence. It helps detect fraud in banking. The people who made this system wanted to make sure it was good at finding fraud and also easy to understand. They used a method called gradient boosting learning and added some techniques to make it explain things in a way that people can understand. This way the system can find fraud accurately. Also give good reasons for its decisions. This is important because it helps banks follow the rules they have to follow. The Explainable Artificial Intelligence system is useful, for banking and it helps with fraud detection.

The system uses a five-stage process to change raw transaction information into predictions that make sense. It starts by getting the data ready which means it gives labels to things and makes sure all the transaction details are on the scale. This includes details about people how they pay and when they make transactions. Then it makes the data more useful by adding details like how far apart things are, when something might be a risk and how much money is involved. The system uses a combination of two methods, called XGBoost and LightGBM to find fraud. It does this by looking at the results from both methods and using a way of training that makes sure it is fair. The system also has a way to explain how it makes its predictions. It uses two methods, called SHAP and LIME to figure out which details are important for each prediction. This helps people who look for fraud and those who have to check that everything is working correctly. The framework makes it easier for them to do their jobs by giving them a list of details, for each prediction.

The ensemble was tested on a sample of 10,000 transactions. This test shows that the ensemble works well. It gets an F1-Score of 0.9438 and an AUC-ROC of 0.9942. These numbers are better than the targets. The framework does what the GDPR Article 22 and EU AI Act say it should do. It makes explanations that can be checked. This helps with using machine learning in places where rules are strict. The ensemble is good at machine learning. It also works well when it is actually being used. This is what is needed for places, with a lot of rules. The. The ensemble meet these needs.

Future work will look at using Apache Kafka to make scoring happen fast in less than 100 milliseconds. This will be for transactions. We will also try to find ways for places to share information about fraud without having to send all the data to one place. This is called learning and it helps keep peoples information private. We want to see how fraudsters behave over time when they make transactions. So we will use something called temporal sequence modeling to look at this. We also need to make sure our system is safe, from people who try to trick it. This is called robustness testing. We will test to see if people can manipulate the explanations that our system gives for why something might be fraud. We will use Apache Kafka for the real-time streaming deployment. This will help us with the transaction scoring.

REFERENCES

- [1] Association of Certified Fraud Examiners (ACFE), “Report to the Nations 2022: Global Study on Occupational Fraud and Abuse,” ACFE, Austin, TX, 2022.
- [2] European Parliament and Council, “Proposal for a Regulation on Artificial Intelligence (AI Act),” COM/2021/206 final, European Commission, Brussels, 2021.
- [3] European Parliament and Council, “Regulation (EU) 2016/679 on the protection of natural persons with regard to the processing of personal data (GDPR),” Official Journal of the European Union, L119, pp. 1–88, 2016.
- [4] PCI Security Standards Council, “Payment Card Industry Data Security Standard (PCI DSS) v4.0,” PCI Security Standards Council, 2022.
- [5] S. M. Lundberg and S.-I. Lee, “A Unified Approach to Interpreting Model Predictions,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 4765–4774.
- [6] M. T. Ribeiro, S. Singh, and C. Guestrin, ““Why Should I Trust You?”: Explaining the Predictions of Any Classifier,” in *Proc. 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 1135–1144.
- [7] T. Chen and C. Guestrin, “XGBoost: A Scalable Tree Boosting System,” in *Proc. 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [8] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, “LightGBM: A Highly Efficient Gradient Boosting Decision Tree,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 3146–3154.
- [9] R. J. Bolton and D. J. Hand, “Statistical Fraud Detection: A Review,” *Statistical Science*, vol. 17, no. 3, pp. 235–255, 2002.
- [10] L. Breiman, “Random Forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [11] F. Pedregosa et al., “Scikit-learn: Machine Learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [12] A. B. Arrieta et al., “Explainable Artificial Intelligence (XAI): Concepts, Taxonomies, Opportunities and Challenges toward Responsible AI,” *Information Fusion*, vol. 58, pp. 82–115, 2020.